

X86 workshop

AsmMasters

Winter 2024



Using scanf & printf



Register	Usage	callee saved
%rax	temporary register; with variable arguments passes information about the number of vector registers used; 1 st return register	No
%rbx	callee-saved register	Yes
%rcx	used to pass 4 th integer argument to functions	No
%rdx	used to pass 3 rd argument to functions; 2 nd return register	No
%rsp	stack pointer	Yes
%rbp	callee-saved register; optionally used as frame pointer	Yes
%rsi	used to pass 2 nd argument to functions	No
%rdi	used to pass 1 st argument to functions	No
%r8	used to pass 5 th argument to functions	No
%r9	used to pass 6 th argument to functions	No
%r10	temporary register, used for passing a function's static chain pointer	No
%r11	temporary register	No
%r12-%r14	callee-saved registers	Yes
%r15	callee-saved register; optionally used as GOT base pointer	Yes
%r16-%r31	temporary registers	No
%xmm0-%xmm1	used to pass and return floating point arguments	No
%xmm2-%xmm7	used to pass floating point arguments	No
%xmm8-%xmm15	temporary registers	No
%xmm16-%xmm31	temporary registers	No
%tmm0-%tmm7	temporary registers	No
%mm0-%mm7	temporary registers	No
%k0-%k7	temporary registers	No
%st0,%st1	temporary registers, used to return long double arguments	No
%st2-%st7	temporary registers	No
%fs	thread pointer	Yes
mxcsr	SSE2 control and status word	partial
x87 SW	x87 status word	No
x87 CW	x87 control word	Yes
tilecfg	Tile control register	No



Classification The size of each argument gets rounded up to eightbytes.¹⁴

The basic types are assigned their natural classes:

- Arguments of types (signed and unsigned) `_Bool`, `char`, `short`, `int`, `long`, `long long`, and pointers are in the INTEGER class.
- Arguments of types `_Float16`, `float`, `double`, `_Decimal32`, `_Decimal64` and `__m64` are in class SSE.
- Arguments of types `__float128`, `_Decimal128` and `__m128` are split into two halves. The least significant ones belong to class SSE, the most significant one to class SSEUP.
- Arguments of type `__m256` are split into four eightbyte chunks. The least significant one belongs to class SSE and all the others to class SSEUP.
- Arguments of type `__m512` are split into eight eightbyte chunks. The least significant one belongs to class SSE and all the others to class SSEUP.



Passing Once arguments are classified, the registers get assigned (in **left-to-right** order) for passing as follows:

1. If the class is **MEMORY**, pass the argument on the stack at an address respecting the arguments alignment (which might be more than its natural alignment).
2. If the class is **INTEGER**, the next available register of the sequence **%rdi, %rsi, %rdx, %rcx, %r8 and %r9** is used¹⁹.
3. If the class is **SSE**, the next available vector register is used, the registers are taken in the order from **%xmm0 to %xmm7**.
4. If the class is **SSEUP**, the eightbyte is passed in the next available eightbyte chunk of the last used vector register.
5. If the class is **X87, X87UP or COMPLEX_X87**, it is passed in memory.

When a value of type `_Bool` is returned or passed in a register or on the stack, bit 0 contains the truth value and bits 1 to 7 shall be zero²⁰.



Returning of Values The returning of values is done according to the following algorithm:

1. Classify the return type with the classification algorithm.
2. If the type has class MEMORY, then the caller provides space for the return value and passes the address of this storage in `%rdi` as if it were the first argument to the function. In effect, this address becomes a “hidden” first argument. This storage must not overlap any data visible to the callee through other names than this argument.

On return `%rax` will contain the address that has been passed in by the caller in `%rdi`.

3. If the class is INTEGER, the next available register of the sequence `%rax, %rdx` is used.
4. If the class is SSE, the next available vector register of the sequence `%xmm0, %xmm1` is used.
5. If the class is SSEUP, the eightbyte is returned in the next available eightbyte chunk of the last used vector register.



The end of the input argument area shall be aligned on a 16 (32 or 64, if `__m256` or `__m512` is passed on stack) byte boundary.¹¹ In other words, the stack needs to be 16 (32 or 64) byte aligned immediately before the call instruction is executed.



Using syscall



1. User-level applications use as integer registers for passing the sequence `%rdi`, `%rsi`, `%rdx`, `%rcx`, `%r8` and `%r9`. The kernel interface uses `%rdi`, `%rsi`, `%rdx`, `%r10`, `%r8` and `%r9`.
2. A system-call is done via the `syscall` instruction. The kernel clobbers registers `%rcx` and `%r11` but preserves all other registers except `%rax`.
3. The number of the syscall has to be passed in register `%rax`.
4. System-calls are limited to six arguments, no argument is passed directly on the stack.
5. Returning from the `syscall`, register `%rax` contains the result of the system-call. A value in the range between -4095 and -1 indicates an error, it is `-errno`.
6. Only values of class INTEGER or class MEMORY are passed to the kernel.



Stack Frame



Position	Contents	Frame
$8n+16$ (%rbp)	memory argument eightbyte n	Previous
	...	
16 (%rbp)	memory argument eightbyte 0	Current
8 (%rbp)	return address	
0 (%rbp)	previous %rbp value	
-8 (%rbp)	unspecified	
	...	
0 (%rsp)	variable size	
-128 (%rsp)	red zone	



Once registers are assigned, the arguments passed in memory are pushed on the stack in reversed (right-to-left²¹) order.



References



- [Guide to x86 Assembly](#)
- [x86 Assembly/X86 Architecture - Wikibooks, open books for an open world](#)
- [ABI](#)
- [x86 and amd64 instruction reference](#)
- [NASM - The Netwide Assembler](#)